

Chapter 23

For Expressions Revisited

Chapter 16 demonstrated that higher-order functions, such as `map`, `flatMap`, and `filter`, provide powerful constructions for dealing with lists. But sometimes the level of abstraction required by these functions makes a program a bit hard to understand.

Here's an example. Say you are given a list of persons, each defined as an instance of a class `Person`. Class `Person` has fields indicating the person's name, whether he or she is male, and his or her children.

Here's the class definition:

```
scala> case class Person(name: String,
                          isMale: Boolean,
                          children: Person*)
```

Here's a list of some sample persons:

```
val lara = Person("Lara", false)
val bob = Person("Bob", true)
val julie = Person("Julie", false, lara, bob)
val persons = List(lara, bob, julie)
```

Now, say you want to find out the names of all pairs of mothers and their children in that list.

Using `map`, `flatMap` and `filter`, you can formulate the following query:

```
scala> persons filter (p => !p.isMale) flatMap (p =>
          (p.children map (c => (p.name, c.name))))
res0: List[(String, String)] = List((Julie,Lara),
          (Julie,Bob))
```

You could optimize this example a bit by using a `withFilter` call instead of `filter`. This would avoid the creation of an intermediate data structure for female persons:

```
scala> persons withFilter (p => !p.isMale) flatMap (p =>
          (p.children map (c => (p.name, c.name))))
res1: List[(String, String)] = List((Julie,Lara),
          (Julie,Bob))
```

These queries do their job, but they are not exactly trivial to write or understand. Is there a simpler way? In fact, there is. Remember the `for` expressions in Section 7.3? Using a `forexpression`, the same example can be written as follows:

```
scala> for (p <- persons; if !p.isMale; c <- p.children)
          yield (p.name, c.name)
res2: List[(String, String)] = List((Julie,Lara),
          (Julie,Bob))
```

The result of this expression is exactly the same as the result of the previous expression. What's more, most readers of the code would likely find the `for` expression much clearer than the previous query, which used the higher-order functions, `map`, `flatMap`, and `withFilter`.

However, the last two queries are not as dissimilar as it might seem. In fact, it turns out that the Scala compiler will translate the second query into the first one. More generally, all foreexpressions that yield a result are translated by the compiler into combinations of invocations of the higher-order methods `map`, `flatMap`, and `withFilter`. All `for` loops without `yield` are translated into a smaller set of higher-order functions: just `withFilter` and `foreach`.

In this chapter, you'll find out first about the precise rules of writing `for` expressions. After that, you'll see how they can make combinatorial problems easier to solve. Finally, you'll learn how `for` expressions are translated, and how as a result, `for` expressions can help you "grow" the Scala language into new application domains.

23.1 FOR EXPRESSIONS

Generally, a `for` expression is of the form:

```
for ( seq ) yield expr
```

Here, `seq` is a sequence of generators, definitions, and filters, with semicolons between successive elements. An example is the `for` expression:

```
for (p <- persons; n = p.name; if (n startsWith "To"))
  yield n
```

This `for` expression contains one generator, one definition, and one filter. As mentioned in Section 7.3 here, you can also enclose the sequence in braces instead of parentheses. Then the semicolons become optional:

```
for {
  p <- persons           // a generator
  n = p.name             // a definition
  if (n startsWith "To") // a filter
} yield n
```

A generator is of the form:

```
pat <- expr
```

The expression `expr` typically returns a list, even though you will see later that this can be generalized. The pattern `pat` gets matched one-by-one against all elements of that list. If the match succeeds, the variables in the pattern get bound to the corresponding parts of the element, just the way it is described in Chapter 15. But if the match fails, no `MatchError` is thrown. Instead, the element is simply discarded from the iteration.

In the most common case, the pattern `pat` is just a variable `x`, as in `x <- expr`. In that case, the variable `x` simply iterates over all elements returned by `expr`.

A definition is of the form:

```
pat = expr
```

This definition binds the pattern `pat` to the value of `expr`, so it has the same effect as a `val` definition:

```
val x = expr
```

The most common case is again where the pattern is a simple variable `x` (e.g., `x = expr`). This defines `x` as a name for the value `expr`.

A filter is of the form:

```
if expr
```

Here, `expr` is an expression of type `Boolean`. The filter drops from the iteration all elements for which `expr` returns false.

Every `for` expression starts with a generator. If there are several generators in a `forexpression`, later generators vary more rapidly than earlier ones. You can verify this easily with the following simple test:

```
scala> for (x <- List(1, 2); y <- List("one", "two"))
  yield (x, y)
res3: List[(Int, String)] =
  List((1,one), (1,two), (2,one), (2,two))
```

23.2 THE N-QUEENS PROBLEM

A particularly suitable application area of `for` expressions are combinatorial puzzles. An example of such a puzzle is the 8-queens problem: Given a standard chess-board, place eight queens such that no queen is in check from any other (a queen can check another piece if they are on the same column, row, or diagonal). To find a solution to this problem, it's actually simpler to generalize it to chess-boards of arbitrary size. Hence, the problem is to place N queens on a chess-board of $N \times N$ squares, where the size N is arbitrary. We'll start numbering cells at one, so the upper-left cell of an $N \times N$ board has coordinate (1, 1) and the lower-right cell has coordinate (N , N).

To solve the N -queens problem, note that you need to place a queen in each row. So you could place queens in successive rows, each time checking that a newly placed queen is not in check from any other queens that have already been placed. In the course of this search, it might happen that a queen that needs to be placed in row k would be in check in all fields of that row from queens in row 1 to $k-1$. In that case, you need to abort that part of the search in order to continue with a different configuration of queens in columns 1 to $k-1$.

An imperative solution to this problem would place queens one by one, moving them around on the board. But it looks difficult to come up with a scheme that really tries all possibilities. A more functional approach represents a solution directly, as a value. A solution consists of a list of coordinates, one for each queen placed on the board. Note, however, that a full solution can not be found in a single step. It needs to be built up gradually, by occupying successive rows with queens.

This suggests a recursive algorithm. Assume you have already generated all solutions of placing k queens on a board of size $N \times N$, where k is less than N . Each such solution can be presented by a list of length k of coordinates (row, column), where both row and column numbers range from 1 to N . It's convenient to treat these partial solution lists as stacks, where the coordinates of the queen in row k come first in the list, followed by the coordinates of the queen in row $k-1$, and so on. The bottom of the stack is the coordinate of the queen placed in the first row of the board. All solutions together are represented as a list of lists, with one element for each solution.

Now, to place the next queen in row $k+1$, generate all possible extensions of each previous solution by one more queen. This yields another list of solution lists, this time of length $k+1$. Continue the process until you have obtained all solutions of the size of the chess-board N .

This algorithmic idea is embodied in function `placeQueens` below:

```
def queens(n: Int): List[List[(Int, Int)]] = {
  def placeQueens(k: Int): List[List[(Int, Int)]] =
    if (k == 0)
      List(List())
    else
      for {
        queens <- placeQueens(k - 1)
        column <- 1 to n
        queen = (k, column)
        if isSafe(queen, queens)
      } yield queen :: queens
  placeQueens(n)
}
```

The outer function `queens` in the program above simply calls `placeQueens` with the size of the board n as its argument. The task of the function application `placeQueens(k)` is to generate all partial solutions of length k in a list. Every element of the list is one solution, represented by a list of length k . So `placeQueens` returns a list of lists.

If the parameter k to `placeQueens` is 0, this means that it needs to generate all solutions of placing zero queens on zero rows. There is only one such solution: place no queen at all. This solution is represented by the empty list. So if k is zero, `placeQueens` returns `List(List())`, a list consisting of a single element that is the empty list. Note that this is quite different from the empty list `List()`.

If `placeQueens` returns `List()`, this means no solutions, instead of a single solution consisting of no placed queens.

In the other case, where k is not zero, all the work of `placeQueens` is done in a `for` expression. The first generator of that `for` expression iterates through all solutions of placing $k - 1$ queens on the board. The second generator iterates through all possible columns on which the k 'th queen might be placed. The third part of the `for` expression defines the newly considered queenposition to be the pair consisting of row k and each produced column. The fourth part of the `forexpression` is a filter which checks with `isSafe` whether the new queen is safe from check by all previous queens (the definition of `isSafe` will be discussed a bit later).

If the new queen is not in check from any other queens, it can form part of a partial solution, so `placeQueens` generates with `queen :: queens` a new solution. If the new queen is not safe from check, the filter returns false, so no solution is generated.

The only remaining bit is the `isSafe` method, which is used to check whether a given queen is in check from any other element in a list of queens. Here is its definition:

```
def isSafe(queen: (Int, Int), queens: List[(Int, Int)]) =
  queens forall (q => !inCheck(queen, q))

def inCheck(q1: (Int, Int), q2: (Int, Int)) =
  q1._1 == q2._1 || // same row
  q1._2 == q2._2 || // same column
  (q1._1 - q2._1).abs == (q1._2 - q2._2).abs // on diagonal
```

The `isSafe` method expresses that a queen is safe with respect to some other queens if it is not in check from any other queen. The `inCheck` method expresses that queens `q1` and `q2` are mutually in check.

It returns true in one of three cases:

1. If the two queens have the same row coordinate,
2. If the two queens have the same column coordinate,
3. If the two queens are on the same diagonal (*i.e.*, the difference between their rows and the difference between their columns are the same).

The first case—that the two queens have the same row coordinate—cannot happen in the application because `placeQueens` already takes care to place each queen in a different row. So you could remove the test without changing the functionality of the program.

23.3 QUERYING WITH FOR EXPRESSIONS

The `for` notation is essentially equivalent to common operations of database query languages. For instance, say you are given a database named `books`, represented as a list of books, where `Book` is defined as follows:

```
case class Book(title: String, authors: String*)
```

Here is a small example database represented as an in-memory list:

```
val books: List[Book] =
  List(
    Book(
      "Structure and Interpretation of Computer Programs",
      "Abelson, Harold", "Sussman, Gerald J."
    ),
    Book(
      "Principles of Compiler Design",
      "Aho, Alfred", "Ullman, Jeffrey"
    ),
    Book(
      "Programming in Modula-2",
      "Wirth, Niklaus"
    ),
  ),
```

```

    Book(
      "Elements of ML Programming",
      "Ullman, Jeffrey"
    ),
    Book(
      "The Java Language Specification", "Gosling, James",
      "Joy, Bill", "Steele, Guy", "Bracha, Gilad"
    )
  )
)

```

To find the titles of all books whose author's last name is "Gosling":

```

scala> for (b <- books; a <- b.authors
           if a.startsWith "Gosling")
       yield b.title
res4: List[String] = List(The Java Language Specification)

```

Or to find the titles of all books that have the string "Program" in their title:

```

scala> for (b <- books if (b.title indexOf "Program") >= 0)
       yield b.title
res5: List[String] = List(Structure and Interpretation of
Computer Programs, Programming in Modula-2, Elements of ML
Programming)

```

Or to find the names of all authors who have written at least two books in the database:

```

scala> for (b1 <- books; b2 <- books if b1 != b2;
           a1 <- b1.authors; a2 <- b2.authors if a1 == a2)
       yield a1
res6: List[String] = List(Ullman, Jeffrey, Ullman, Jeffrey)

```

The last solution is still not perfect because authors will appear several times in the list of results. You still need to remove duplicate authors from result lists. This can be achieved with the following function:

```

scala> def removeDuplicates[A](xs: List[A]): List[A] = {
           if (xs.isEmpty) xs
           else
             xs.head :: removeDuplicates(
                           xs.tail filter (x => x != xs.head)
                         )
         }
removeDuplicates: [A](xs: List[A])List[A]

scala> removeDuplicates(res6)
res7: List[String] = List(Ullman, Jeffrey)

```

It's worth noting that the last expression in method `removeDuplicates` can be equivalently expressed using a `for` expression:

```

xs.head :: removeDuplicates(
  for (x <- xs.tail if x != xs.head) yield x
)

```

23.4 TRANSLATION OF FOR EXPRESSIONS

Every for expression can be expressed in terms of the three higher-order functions `map`, `flatMap`, and `withFilter`. This section describes the translation scheme, which is also used by the Scala compiler.

Translating for expressions with one generator

First, assume you have a simple for expression:

```
for (null <- expr_1) yield expr_2
```

where `x` is a variable. Such an expression is translated to:

```
expr_1.map(null => expr_2)
```

Translating for expressions starting with a generator and a filter

Now, consider for expressions that combine a leading generator with some other elements.

A for expression of the form:

```
for (null <- expr_1 if expr_2) yield expr_3
```

is translated to:

```
for (null <- expr_1 withFilter (null => expr_2)) yield expr_3
```

This translation gives another for expression that is shorter by one element than the original, because an `if` element is transformed into an application of `withFilter` on the first generator expression. The translation then continues with this second expression, so in the end you obtain:

```
expr_1 withFilter (null => expr_2) map (null => expr_3)
```

The same translation scheme also applies if there are further elements following the filter. If `seq` is an arbitrary sequence of generators, definitions, and filters, then:

```
for (null <- expr_1 if expr_2; seq) yield expr_3
```

is translated to:

```
for (null <- expr_1 withFilter expr_2; seq) yield expr_3
```

Then translation continues with the second expression, which is again shorter by one element than the original one.

Translating for expressions starting with two generators

The next case handles for expressions that start with two generators, as in:

```
for (null <- expr_1; null <- expr_2; seq) yield expr_3
```

Again, assume that `seq` is an arbitrary sequence of generators, definitions, and filters. In fact, `seq` might also be empty, and in that case there would not be a semicolon after `expr_2`. The translation scheme stays the same in each case. The `for` expression above is translated to an application of `flatMap`:

```
expr_1.flatMap(null => for (null <- expr_2; seq) yield expr_3)
```

This time, there is another `for` expression in the function value passed to `flatMap`. That `for` expression (which is again simpler by one element than the original) is in turn translated with the same rules.

The three translation schemes given so far are sufficient to translate all `for` expressions that contain just generators and filters, and where generators bind only simple variables. Take, for instance, the query, "find all authors who have published at least two books," from Section 23.3:

```
for (b1 <- books; b2 <- books if b1 != b2;  
    a1 <- b1.authors; a2 <- b2.authors if a1 == a2)  
yield a1
```

This query translates to the following `map/flatMap/filter` combination:

```
books flatMap (b1 =>  
  books withFilter (b2 => b1 != b2) flatMap (b2 =>  
    b1.authors flatMap (a1 =>  
      b2.authors withFilter (a2 => a1 == a2) map (a2 =>  
        a1))))
```

The translation scheme presented so far does not yet handle generators that bind whole patterns instead of simple variables. It also does not yet cover definitions. These two aspects will be explained in the next two sub-sections.

Translating patterns in generators

The translation scheme becomes more complicated if the left hand side of generator is a pattern, `pat`, other than a simple variable. The case where the `for` expression binds a tuple of variables is still relatively easy to handle. In that case, almost the same scheme as for single variables applies.

A `for` expression of the form:

```
for ((null, ..., null) <- expr_1) yield expr_2
```

translates to:

```
expr_1.map { case (null, ..., null) => expr_2 }
```

Things become a bit more involved if the left hand side of the generator is an arbitrary pattern `pat` instead of a single variable or a tuple.

In this case:

```
for (pat <- expr_1) yield expr_2
```

translates to:


```

expr_1 withFilter {
  case pat => true
  case _ => false
} map {
  case pat => expr_2
}

```

That is, the generated items are first filtered and only those that match `pat` are mapped. Therefore, it's guaranteed that a pattern-matching generator will never throw a `MatchError`.

The scheme here only treated the case where the for expression contains a single pattern-matching generator. Analogous rules apply if the for expression contains other generators, filters or definitions. Because these additional rules don't add much new insight, they are omitted from discussion here. If you are interested, you can look them up in the *Scala Language Specification* [Ode11].

Translating definitions

The last missing situation is where a for expression contains embedded definitions. Here's a typical case:

```
for (null <- expr_1; null = expr_2; seq) yield expr_3
```

Assume again that `seq` is a (possibly empty) sequence of generators, definitions, and filters. This expression is translated to this one:

```
for ((null, null) <- for (null <- expr_1) yield (null, expr_2); seq)
yield expr_3
```

So you see that `expr_2` is evaluated each time there is a new `x` value being generated. This re-evaluation is necessary because `expr_2` might refer to `x` and so needs to be re-evaluated for changing values of `x`. For you as a programmer, the conclusion is that it's probably not a good idea to have definitions embedded in for expressions that do not refer to variables bound by some preceding generator, because re-evaluating such expressions would be wasteful. For instance, instead of:

```
for (x <- 1 to 1000; y = expensiveComputationNotInvolvingX)
yield x * y
```

it's usually better to write:

```
val y = expensiveComputationNotInvolvingX
for (x <- 1 to 1000) yield x * y
```

Translating for loops

The previous subsections showed how for expressions that contain a `yield` are translated. What about for loops that simply perform a side effect without returning anything? Their translation is similar, but simpler than for expressions. In principle, wherever the previous translation scheme used a `map` or a `flatMap` in the translation, the translation scheme for for loops uses just a `foreach`.

For instance, the expression:

```
for (null <- expr_1) body
```

translates to:

```
expr_1 foreach (null => body)
```

A larger example is the expression:

```
for (null <- expr_1; if expr_2; null <- expr_3) body
```

This expression translates to:

```
expr_1 withFilter (null => expr_2) foreach (null =>
  expr_3 foreach (null => body))
```

For example, the following expression sums up all elements of a matrix represented as a list of lists:

```
var sum = 0
for (xs <- xss; x <- xs) sum += x
```

This loop is translated into two nested foreach applications:

```
var sum = 0
xss foreach (xs =>
  xs foreach (x =>
    sum += x))
```

23.5 GOING THE OTHER WAY

The previous section showed that for expressions can be translated into applications of the higher-order functions `map`, `flatMap`, and `withFilter`. In fact, you could equally go the other way: Every application of a `map`, `flatMap`, or `filter` can be represented as a `for` expression.

Here are implementations of the three methods in terms of `for` expressions. The methods are contained in an object `Demo` to distinguish them from the standard operations on `Lists`. To be concrete, the three functions all take a `List` as parameter, but the translation scheme would work just as well with other collection types:

```
object Demo {
  def map[A, B](xs: List[A], f: A => B): List[B] =
    for (x <- xs) yield f(x)

  def flatMap[A, B](xs: List[A], f: A => List[B]): List[B] =
    for (x <- xs; y <- f(x)) yield y

  def filter[A](xs: List[A], p: A => Boolean): List[A] =
    for (x <- xs if p(x)) yield x
}
```

Not surprisingly, the translation of the `for` expression used in the body of `Demo.map` will produce a call to `map` in class `List`. Similarly, `Demo.flatMap` and `Demo.filter` translate to `flatMap` and `withFilter` in class `List`. So this little demonstration shows that `for` expressions really are equivalent in their expressiveness to applications of the three functions `map`, `flatMap`, and `withFilter`.

23.6 GENERALIZING FOR

Because the translation of for expressions only relies on the presence of methods `map`, `flatMap`, and `withFilter`, it is possible to apply the for notation to a large class of data types.

You have already seen for expressions over lists and arrays. These are supported because lists, as well as arrays, define operations `map`, `flatMap`, and `withFilter`. Because they define a `foreach` method as well, for loops over these data types are also possible.

Besides lists and arrays, there are many other types in the Scala standard library that support the same four methods and therefore allow for expressions. Examples are ranges, iterators, streams, and all implementations of sets. It's also perfectly possible for your own data types to support for expressions by defining the necessary methods. To support the full range of forexpressions and for loops, you need to define `map`, `flatMap`, `withFilter`, and `foreach` as methods of your data type. But it's also possible to define a subset of these methods, and thereby support a subset of all possible for expressions or loops.

Here are the precise rules:

- If your type defines just `map`, it allows for expressions consisting of a single generator.
- If it defines `flatMap` as well as `map`, it allows for expressions consisting of several generators.
- If it defines `foreach`, it allows for loops (both with single and multiple generators).
- If it defines `withFilter`, it allows for filter expressions starting with an `if` in the forexpression.

The translation of for expressions happens before type checking. This allows for maximum flexibility because the only requirement is that the result of expanding a for expression type checks. Scala defines no typing rules for the for expressions themselves, and does not require that methods `map`, `flatMap`, `withFilter`, or `foreach` have any particular type signatures.

Nevertheless, there is a typical setup that captures the most common intention of the higher order methods to which for expressions translate. Say you have a parameterized class, `C`, which typically would stand for some sort of collection. Then it's quite natural to pick the following type signatures for `map`, `flatMap`, `withFilter`, and `foreach`:

```
abstract class C[A] {  
  def map[B](f: A => B): C[B]  
  def flatMap[B](f: A => C[B]): C[B]  
  def withFilter(p: A => Boolean): C[A]  
  def foreach(b: A => Unit): Unit  
}
```

That is, the `map` function takes a function from the collection's element type `A` to some other type `B`. It produces a new collection of the same kind `C`, but with `B` as the element type. The `flatMap` method takes a function `f` from `A` to some `C`-collection of `B`s and produces a `C`-collection of `B`s.

The `withFilter` method takes a predicate function from the collection's element type `A` to `Boolean`. It produces a collection of the same type as the one on which it is invoked. Finally, the `foreach` method takes a function from `A` to `Unit` and produces a `Unit` result:

In class `C` above, the `withFilter` method produces a new collection of the same class. That means that every invocation of `withFilter` creates a new `C` object, just the same as `filter` would work. Now, in the translation of `for` expressions, any calls to `withFilter` are always followed by calls to one of the other three methods. Therefore, the object created by `withFilter` will be taken apart by one of the other methods immediately afterwards. If objects of class `C` are large (think long sequences), you might want to avoid the creation of such an intermediate object. A standard technique is to let `withFilter` return not a `C` object but just a wrapper object that "remembers" that elements need to be filtered before being processed further.

Concentrating on just the first three functions of class `C`, the following facts are noteworthy. In functional programming, there's a general concept called a monad, which can explain a large number of types with computations, ranging from collections, to computations with state and I/O, backtracking computations, and transactions, to name a few. You can formulate functions `map`, `flatMap`, and `withFilter` on a monad, and, if you do, they end up having exactly the types given here.

Furthermore, you can characterize every monad by `map`, `flatMap`, and `withFilter`, plus a "unit" constructor that produces a monad from an element value. In an object-oriented language, this "unit" constructor is simply an instance constructor or a factory method. Therefore, `map`, `flatMap`, and `withFilter` can be seen as an object-oriented version of the functional concept of monad. Because `for` expressions are equivalent to applications of these three methods, they can be seen as syntax for monads.

All this suggests that the concept of `for` expression is more general than just iteration over a collection, and indeed it is. For instance, `for` expressions also play an important role in asynchronous I/O, or as an alternative notation for optional values. Watch out in the Scala libraries for occurrences of `map`, `flatMap`, and `withFilter`—when they are present, `for` expressions suggest themselves as a concise way of manipulating elements of the type.

23.7 CONCLUSION

In this chapter, you were given a peek under the hood of `for` expressions and `for` loops. You learned that they translate into applications of a standard set of higher-order methods. As a result, you saw that `for` expressions are really much more general than mere iterations over collections, and that you can design your own classes to support them.