

## Chapter 25

# The Architecture of Scala Collections

This chapter describes the architecture of the Scala collections framework in detail. Continuing the theme of Chapter 24, you will find out more about the internal workings of the framework. You will also learn how this architecture helps you define your own collections in a few lines of code, while reusing the overwhelming part of collection functionality from the framework.

Chapter 24 enumerated a large number of collection operations, which exist uniformly on many different collection implementations. Implementing every collection operation anew for every collection type would lead to an enormous amount of code, most of which would be copied from somewhere else. Such code duplication could lead to inconsistencies over time, when an operation is added or modified in one part of the collection library but not in others. The principal design objective of the new collections framework was to avoid any duplication, defining every operation in as few places as possible.<sup>[1]</sup> The design approach was to implement most operations in collection "templates" that can be flexibly inherited from individual base classes and implementations. In this chapter, we will examine these templates, and other classes and traits that constitute the "building blocks" of the framework, as well as the construction principles they support.

### 25.1 BUILDERS

Almost all collection operations are implemented in terms of traversals and builders. Traversals are handled by Traversable's `foreach` method, and building new collections is handled by instances of class `Builder`. Listing 25.1 presents a slightly abbreviated outline of this class.

```
package scala.collection.generic

class Builder[-Elem, +To] {
  def +=(elem: Elem): this.type
  def result(): To
  def clear()
  def mapResult[NewTo](f: To => NewTo): Builder[Elem, NewTo]
    = ...
}
```

**Listing 25.1 - An outline of the Builder class.**

You can add an element `x` to a builder `b` with `b += x`. There's also syntax to add more than one element at once: For instance, `b += (x, y)` and `b ++= xs` work as for buffers. (In fact, buffers are an enriched version of builders.) The `result()` method returns a collection from a builder. The state of the builder is undefined after taking its result, but it can be reset into a new empty state using `clear()`. Builders are generic in both the element type, `Elem`, and in the type, `To`, of collections they return.

Often, a builder can refer to some other builder for assembling the elements of a collection, but then would like to transform the result of the other builder—for example, to give it a different type. This task is simplified by method `mapResult` in class `Builder`. Suppose for instance you have an array

buffer `buf`. Array buffers are builders for themselves, so taking `result()` of an array buffer will return the same buffer. If you want to use this buffer to produce a builder that builds arrays, you can use `mapResult`:

```
scala> val buf = new ArrayBuffer[Int]
buf: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer()

scala> val bldr = buf mapResult (_.toArray)
bldr: scala.collection.mutable.Builder[Int,Array[Int]]
      = ArrayBuffer()
```

The result value, `bldr`, is a builder that uses the array buffer, `buf`, to collect elements. When a result is demanded from `bldr`, the result of `buf` is computed, which yields the array buffer `buf` itself. This array buffer is then mapped with `_.toArray` to an array. So the end result is that `bldr` is a builder for arrays.

## 25.2 FACTORING OUT COMMON OPERATIONS

The main design objectives of the collection library redesign were to have, at the same time, natural types and maximal sharing of implementation code. In particular, Scala's collections follow the "same-result-type" principle: Wherever possible, a transformation method on a collection will yield a collection of the same type. For instance, the filter operation should yield, on every collection type, an instance of the same collection type. Applying filter on a `List` should give a `List`; applying it on a `Map` should give a `Map`; and so on. In the rest of this section, you will find out how this is achieved.

### THE FAST TRACK

The material in this section is a bit more dense than usual and might require some time to absorb. If you want to move ahead quickly, you could skip the remainder of this section and move on to Section 25.3 here where you will learn from concrete examples how to integrate your own collection classes in the framework.

The Scala collection library avoids code duplication and achieves the "same-result-type" principle by using generic builders and traversals over collections in so-called *implementation traits*. These traits are named with a `Like` suffix; for instance, `IndexedSeqLike` is the implementation trait for `IndexedSeq`, and similarly, `TraversableLike` is the implementation trait for `Traversable`. Collection classes such as `Traversable` or `IndexedSeq` inherit all their concrete method implementations from these traits. Implementation traits have two type parameters instead of one for normal collections. They parameterize not only over the collection's element type, but also over the collection's representation type (*i.e.*, the type of the underlying collection), such as `Seq[I]` or `List[T]`.

For instance, here is the header of trait `TraversableLike`:

```
trait TraversableLike[+Elem, +Repr] { ... }
```

The type parameter, `Elem`, stands for the element type of the traversable whereas the type parameter `Repr` stands for its representation. There are no constraints on `Repr`. In particular `Repr` might be instantiated to a type that is itself not a subtype of `Traversable`. That way, classes outside the

collections hierarchy, such as `String` and `Array`, can still make use of all operations defined in a collection implementation trait.

```
package scala.collection

trait TraversableLike[+Elem, +Repr] {
  def newBuilder: Builder[Elem, Repr] // deferred
  def foreach[U](f: Elem => U)         // deferred
  ...
  def filter(p: Elem => Boolean): Repr = {
    val b = newBuilder
    foreach { elem => if (p(elem)) b += elem }
    b.result
  }
}
```

#### Listing 25.2 - Implementation of `filter` in `TraversableLike`.

Taking `filter` as an example, this operation is defined once for all collection classes in the trait `TraversableLike`. An outline of the relevant code is shown in Listing 25.2. The trait declares two abstract methods, `newBuilder` and `foreach`, which are implemented in concrete collection classes. The `filter` operation is implemented in the same way for all collections using these methods. It first constructs a new builder for the representation type `Repr`, using `newBuilder`. It then traverses all elements of the current collection, using `foreach`. If an element `x` satisfies the given predicate `p`—*i.e.*, `p(x)` is true—it is added with the builder. Finally, the elements collected in the builder are returned as an instance of the `Repr` collection type by calling the builder's `result` method.

The `map` operation on collections is a bit more complicated. For instance, if `f` is a function from `String` to `Int`, and `xs` is a `List[String]`, then `xs map f` should give a `List[Int]`. Likewise, if `ys` is an `Array[String]`, then `ys map f` should give a `Array[Int]`. But how do you achieve that without duplicating the definition of the `map` method in lists and arrays?

The `newBuilder/foreach` framework shown in Listing 25.2 is not sufficient for this because it only allows creation of new instances of the same collection type, whereas `map` needs an instance of the same collection type constructor but possibly with a different element type. What's more, even the result type constructor of a function like `map` might depend, in non-trivial ways, on the other argument types. Here is an example:

```
scala> import collection.immutable.BitSet
import collection.immutable.BitSet

scala> val bits = BitSet(1, 2, 3)
bits: scala.collection.immutable.BitSet = BitSet(1, 2, 3)

scala> bits map (_ * 2)
res13: scala.collection.immutable.BitSet = BitSet(2, 4, 6)

scala> bits map (_.toFloat)
res14: scala.collection.immutable.Set[Float] =
  Set(1.0, 2.0, 3.0)
```

If you map the doubling function `_ * 2` over a bit set you obtain another bit set. However, if you map the function `(_.toFloat)` over the same bit set, the result is a general `Set[Float]`. Of course, it can't be a bit set because bit sets contain `Int`s, not `Float`s.

Note that `map`'s result type depends on the type of function that's passed to it. If the result type of that function argument is again an `Int`, the result of `map` is a `BitSet`. But if the result type of the function argument is something else, the result of `map` is just a `Set`. You'll find out soon how this type-flexibility is achieved in Scala.

The problem with `BitSet` is not an isolated case. Here are two more interactions with the interpreter that both map a function over a map:

```
scala> Map("a" -> 1, "b" -> 2) map { case (x, y) => (y, x) }
res3: scala.collection.immutable.Map[Int,java.lang.String] =
  Map(1 -> a, 2 -> b)

scala> Map("a" -> 1, "b" -> 2) map { case (x, y) => y }
res4: scala.collection.immutable.Iterable[Int] =
  List(1, 2)
```

The first function swaps two arguments of a key/value pair. The result of mapping this function is again a map, but now going in the other direction. In fact, the first expression yields the inverse of the original map, provided it is invertible. The second function, however, maps the key/value pair to an integer, namely its value component. In that case, we cannot form a `Map` from the results, but we still can form an `Iterable`, a supertrait of `Map`.

You might ask, Why not restrict `map` so that it can always return the same kind of collection? For instance, on bit sets `map` could accept only `Int`-to-`Int` functions and on maps it could only accept pair-to-pair functions. Not only are such restrictions undesirable from an object-oriented modeling point of view, they are illegal because they would violate the Liskov Substitution Principle:

A `Map` is an `Iterable`. So every operation that's legal on an `Iterable` must also be legal on a `Map`.

Scala solves this problem instead with overloading: Not the simple form of overloading inherited by Java (that would not be flexible enough), but the more systematic form of overloading that's provided by implicit parameters.

```
def map[B, That](f: Elem => B)
  (implicit bf: CanBuildFrom[Repr, B, That]): That = {
  val b = bf(this)
  for (x <- this) b += f(x)
  b.result
}
```

#### **Listing 25.3 - Implementation of map in TraversableLike.**

Listing 25.3 shows trait `TraversableLike`'s implementation of `map`. It's quite similar to the implementation of `filter` shown in Listing 25.2. The principal difference is that where `filter` used the `newBuilder` method, which is abstract in class `TraversableLike`, `map` uses a builder factory that's passed as an additional implicit parameter of type `CanBuildFrom`.

```

package scala.collection.generic

trait CanBuildFrom[-From, -Elem, +To] {
  // Creates a new builder
  def apply(from: From): Builder[Elem, To]
}

```

#### Listing 25.4 - The CanBuildFrom trait.

Listing 25.4 shows the definition of the trait `CanBuildFrom`, which represents builder factories. It has three type parameters: `Elem` indicates the element type of the collection to be built, `To` indicates the type of collection to build, and `From` indicates the type for which this builder factory applies. By defining the right implicit definitions of builder factories, you can tailor the right typing behavior as needed.

Take class `BitSet` as an example. Its companion object would contain a builder factory of type `CanBuildFrom[BitSet, Int, BitSet]`. This means that when operating on a `BitSet` you can construct another `BitSet`, provided the type of the collection to build is `Int`. If this is not the case, you can always fall back to a different implicit builder factory, this time implemented in `immutable.Set`'s companion object. The type of this more general builder factory, where `A` is a generic type parameter, is:

```
CanBuildFrom[Set[_], A, Set[A]]
```

This means that when operating on an arbitrary `Set`, expressed by the wildcard type `Set[_]`, you can build a `Set` again no matter what the element type `A` is. Given these two implicit instances of `CanBuildFrom`, you can then rely on Scala's rules for implicit resolution to pick the one that's appropriate and maximally specific.

So implicit resolution provides the correct static types for tricky collection operations, such as `map`. But what about the dynamic types? Specifically, say you have a list value that has `Iterable` as its static type, and you map some function over that value:

```

scala> val xs: Iterable[Int] = List(1, 2, 3)
xs: Iterable[Int] = List(1, 2, 3)

scala> val ys = xs map (x => x * x)
ys: Iterable[Int] = List(1, 4, 9)

```

The static type of `ys` above is `Iterable`, as expected. But its dynamic type is (and should be) still `List`! This behavior is achieved by one more indirection. The `apply` method in `CanBuildFrom` is passed the source collection as argument. Most builder factories for generic traversables (in fact all except builder factories for leaf classes) forward the call to a method `genericBuilder` of a collection. The `genericBuilder` method in turn calls the builder that belongs to the collection in which it is defined. So Scala uses static implicit resolution to resolve constraints on the types of `map`, and virtual dispatch to pick the best dynamic type that corresponds to these constraints.

## 25.3 INTEGRATING NEW COLLECTIONS

What needs to be done if you want to integrate a new collection class, so that it can profit from all predefined operations at the right types? In this section we'll show you two examples that do this.

## Integrating sequences

Say you want to create a new sequence type for RNA strands, which are sequences of bases A (adenine), T (thymine), G (guanine), and U (uracil). The definitions for bases are easily set up as shown in Listing 25.5.

Every base is defined as a case object that inherits from a common abstract class `Base`. The `Base` class has a companion object that defines two functions that map between bases and the integers 0 to 3. You can see in the examples two different ways to use collections to implement these functions.

The `toInt` function is implemented as a `Map` from `Base` values to integers. The reverse function, `fromInt`, is implemented as an array. This makes use of the fact that both maps and arrays are functions because they inherit from the `Function1` trait.

The next task is to define a class for strands of RNA. Conceptually, a strand of RNA is simply a `Seq[Base]`. However, RNA strands can get quite long, so it makes sense to invest some work in a compact representation. Because there are only four bases, a base can be identified with two bits, and you can therefore store sixteen bases as two-bit values in an integer. The idea, then, is to construct a specialized subclass of `Seq[Base]`, which uses this packed representation.

Listing 25.6 presents the first version of this class; it will be refined later. The class `RNA1` has a constructor that takes an array of `Ints` as its first argument. This array contains the packed RNA data, with sixteen bases in each element, except for the last array element, which might be partially filled. The second argument, `length`, specifies the total number of bases on the array (and in the sequence). Class `RNA1` extends `IndexedSeq[Base]`. Trait `IndexedSeq`, which comes from package `scala.collection.immutable`, defines two abstract methods, `length` and `apply`.

```
abstract class Base
case object A extends Base
case object T extends Base
case object G extends Base
case object U extends Base

object Base {
  val fromInt: Int => Base = Array(A, T, G, U)
  val toInt: Base => Int = Map(A -> 0, T -> 1, G -> 2, U -> 3)
}
```

### Listing 25.5 - RNA Bases.

These need to be implemented in concrete subclasses. Class `RNA1` implements `length` automatically by defining a parametric field (described in Section 10.6) of the same name. It implements the indexing method `apply` with the code given in Listing 25.6. Essentially, `apply` first extracts an integer value from the `groups` array, then extracts the correct two-bit number from that integer using right shift (`>>`) and mask (`&`). The private constants `S`, `N`, and `M` come from the `RNA1` companion object. `S` specifies the size of each packet (*i.e.*, two); `N` specifies the number of two-bit packets per integer; and `M` is a bit mask that isolates the lowest `S` bits in a word.

Note that the constructor of class `RNA1` is private. This means that clients cannot create `RNA1` sequences by calling `new`, which makes sense, because it hides the representation of `RNA1` sequences in terms of packed arrays from the user. If clients cannot see what the representation details of RNA sequences are, it becomes possible to change these representation details at any point in the future without affecting client code.

In other words, this design achieves a good decoupling of the interface of RNA sequences and its implementation. However, if constructing an RNA sequence with `new` is impossible, there must be some other way to create new RNA sequences, or else the whole class would be rather useless. There are two alternatives for RNA sequence creation, both provided by the `RNA1` companion object. The first way is method `fromSeq`, which converts a given sequence of bases (*i.e.*, a value of type `Seq[Base]`) into an instance of class `RNA1`. The `fromSeq` method does this by packing all the bases contained in its argument sequence into an array, then calling `RNA1`'s private constructor with that array and the length of the original sequence as arguments. This makes use of the fact that a private constructor of a class is visible in the class's companion object.

```
import collection.IndexedSeqLike
import collection.mutable.{Builder, ArrayBuffer}
import collection.generic.CanBuildFrom

final class RNA1 private (val groups: Array[Int],
    val length: Int) extends IndexedSeq[Base] {

    import RNA1._

    def apply(idx: Int): Base = {
        if (idx < 0 || length <= idx)
            throw new IndexOutOfBoundsException
        Base.fromInt(groups(idx / N) >> (idx % N * S) & M)
    }
}

object RNA1 {

    // Number of bits necessary to represent group
    private val S = 2

    // Number of groups that fit in an Int
    private val N = 32 / S

    // Bitmask to isolate a group
    private val M = (1 << S) - 1

    def fromSeq(buf: Seq[Base]): RNA1 = {
        val groups = new Array[Int]((buf.length + N - 1) / N)
        for (i <- 0 until buf.length)
            groups(i / N) |= Base.toInt(buf(i)) << (i % N * S)
        new RNA1(groups, buf.length)
    }

    def apply(bases: Base*) = fromSeq(bases)
}
```

#### Listing 25.6 - RNA strands class, first version.

The second way to create an RNA1 value is provided by the apply method in the RNA1 object. It takes a variable number of Base arguments and simply forwards them as a sequence to fromSeq.

Here are the two creation schemes in action:

```
scala> val xs = List(A, G, T, A)
xs: List[Product with Base] = List(A, G, T, A)

scala> RNA1.fromSeq(xs)
res1: RNA1 = RNA1(A, G, T, A)

scala> val rna1 = RNA1(A, U, G, G, T)
rna1: RNA1 = RNA1(A, U, G, G, T)
```

#### Adapting the result type of RNA methods

Here are some more interactions with the RNA1 abstraction:

```
scala> rna1.length
res2: Int = 5

scala> rna1.last
res3: Base = T

scala> rna1.take(3)
res4: IndexedSeq[Base] = Vector(A, U, G)
```

The first two results are as expected, but the last result of taking the first three elements of rna1 might not be. In fact, you see an IndexedSeq[Base] as static result type and a Vector as the dynamic type of the result value. You might have expected to see an RNA1 value instead. But this is not possible because all that was done in Listing 25.6 was make RNA1 extend IndexedSeq. Class IndexedSeq, on the other hand, has a take method that returns an IndexedSeq, and that's implemented in terms of IndexedSeq's default implementation, Vector.

Now that you understand why things are the way they are, the next question should be what needs to be done to change them? One way to do this would be to override the take method in class RNA1, maybe like this:

```
def take(count: Int): RNA1 = RNA1.fromSeq(super.take(count))
```

This would do the job for take. But what about drop, or filter, or init? In fact there are over fifty methods on sequences that return again a sequence. For consistency, all of these would have to be overridden. This looks less and less like an attractive option.

Fortunately, there is a much easier way to achieve the same effect. The RNA class needs to inherit not only from IndexedSeq, but also from its implementation trait IndexedSeqLike. This is shown in Listing 25.7. The new implementation differs from the previous one in only two aspects. First, class RNA2 now also extends from IndexedSeqLike[Base, RNA2]. The IndexedSeqLike trait implements all concrete methods of IndexedSeq in an extensible way.



For instance, the return type of methods like `take`, `drop`, `filter` or `init` is the second type parameter passed to class `IndexedSeqLike` (i.e., `RNA2` in Listing 25.7). To do this, `IndexedSeqLike` bases itself on the `newBuilder` abstraction, which creates a builder of the right kind. Subclasses of trait `IndexedSeqLike` have to override `newBuilder` to return collections of their own kind. In class `RNA2`, the `newBuilder` method returns a builder of type `Builder[Base, RNA2]`. To construct this builder, it first creates an `ArrayBuffer`, which itself is a `Builder[Base, ArrayBuffer]`. It then transforms the `ArrayBuffer` builder by calling its `mapResult` method to an `RNA2` builder. The `mapResult` method expects a transformation function from `ArrayBuffer` to `RNA2` as its parameter. The function given is simply `RNA2.fromSeq`, which converts an arbitrary base sequence to an `RNA2` value (recall that an array buffer is a kind of sequence, so `RNA2.fromSeq` can be applied to it).

```
final class RNA2 private (
  val groups: Array[Int],
  val length: Int
) extends IndexedSeq[Base] with IndexedSeqLike[Base, RNA2] {

  import RNA2._

  override def newBuilder: Builder[Base, RNA2] =
    new ArrayBuffer[Base] mapResult fromSeq

  def apply(idx: Int): Base = // as before
}
```

**Listing 25.7 - RNA strands class, second version.**

If you had left out the `newBuilder` definition, you would have gotten an error message like the following:

```
RNA2.scala:5: error: overriding method newBuilder in trait
TraversableLike of type => scala.collection.mutable.Builder[Base, RNA2];
method newBuilder in trait GenericTraversableTemplate of type
=> scala.collection.mutable.Builder[Base, IndexedSeq[Base]] has
incompatible type
class RNA2 private (val groups: Array[Int], val length: Int)
^
one error found
```

The error message is quite long and complicated, which reflects the intricate way the collection libraries are put together. It's best to ignore the information about where the methods come from, because in this case it detracts more than it helps. What remains is that a method `newBuilder` with result type `Builder[Base, RNA2]` needed to be defined, but a method `newBuilder` with result type `Builder[Base, IndexedSeq[Base]]` was found. The latter does not override the former.

The first method, whose result type is `Builder[Base, RNA2]`, is an abstract method that got instantiated at this type in Listing 25.7 by passing the `RNA2` type parameter to `IndexedSeqLike`. The second method, of result type `Builder[Base, IndexedSeq[Base]]`, is what's provided by the inherited `IndexedSeq` class. In other words, the `RNA2` class is invalid without a definition of `newBuilder` with the first result type.

With the refined implementation of the RNA class in Listing 25.7, methods like `take`, `drop`, `orfilter` work now as expected:

```
scala> val rna2 = RNA2(A, U, G, G, T)
rna2: RNA2 = RNA2(A, U, G, G, T)

scala> rna2 take 3
res5: RNA2 = RNA2(A, U, G)

scala> rna2 filter (U !=)
res6: RNA2 = RNA2(A, G, G, T)
```

### Dealing with map and friends

There is another class of methods in collections that we haven't dealt with yet. These methods do not always return the collection type exactly. They might return the same kind of collection, but with a different element type. The classical example of this is the `map` method. Ifs is a `Seq[Int]`, and `f` is a function from `Int` to `String`, then `s.map(f)` would return a `Seq[String]`. So the element type changes between the receiver and the result, but the kind of collection stays the same.

There are a number of other methods that behave like `map`. For some of them you would expect this (*e.g.*, `flatMap`, `collect`), but for others you might not. For instance, the `append` method, `++`, also might return a result whose type differs from that of its arguments—appending a list of `String` to a list of `Int` would give a list of `Any`. How should these methods be adapted to RNA strands? Ideally we'd expect that mapping bases to bases over an RNA strand would yield again an RNA strand:

```
scala> val rna = RNA(A, U, G, G, T)
rna: RNA = RNA(A, U, G, G, T)

scala> rna map { case A => T case b => b }
res7: RNA = RNA(T, U, G, G, T)
```

Likewise, appending two RNA strands with `++` should yield again another RNA strand:

```
scala> rna ++ rna
res8: RNA = RNA(A, U, G, G, T, A, U, G, G, T)
```

On the other hand, mapping bases to some other type over an RNA strand cannot yield another RNA strand because the new elements have the wrong type. It has to yield a sequence instead. In the same vein appending elements that are not of type `Base` to an RNA strand can yield a general sequence, but it cannot yield another RNA strand.

```
scala> rna map Base.toInt
res2: IndexedSeq[Int] = Vector(0, 3, 2, 2, 1)

scala> rna ++ List("missing", "data")
res3: IndexedSeq[java.lang.Object] =
  Vector(A, U, G, G, T, missing, data)
```

This is what you'd expect in the ideal case. But this is not what the RNA2 class as given in Listing 25.7 provides. In fact, if you ran the first two examples above with instances of this class you would obtain:

```
scala> val rna2 = RNA2(A, U, G, G, T)
rna2: RNA2 = RNA2(A, U, G, G, T)

scala> rna2 map { case A => T case b => b }
res0: IndexedSeq[Base] = Vector(T, U, G, G, T)

scala> rna2 ++ rna2
res1: IndexedSeq[Base] = Vector(A, U, G, G, T, A, U, G, G, T)
```

So the result of map and ++ is never an RNA strand, even if the element type of the generated collection is a Base. To see how to do better, it pays to have a close look at the signature of the map method (or of ++, which has a similar signature). The map method is originally defined in class scala.collection.TraversableLike with the following signature:

```
def map[B, That](f: Elem => B)
  (implicit cbf: CanBuildFrom[Repr, B, That]): That
```

Here Elem is the type of elements of the collection, and Repr is the type of the collection itself; that is, the second type parameter that gets passed to implementation classes such as TraversableLike and IndexedSeqLike. The map method takes two more type parameters, B and That. The B parameter stands for the result type of the mapping function, which is also the element type of the new collection. The That appears as the result type of map, so it represents the type of the new collection that gets created.

How is the That type determined? It is linked to the other types by an implicit parameter cbf, of type CanBuildFrom[Repr, B, That]. These CanBuildFrom implicits are defined by the individual collection classes. In essence, an implicit value of type CanBuildFrom[From, Elem, To] says: "Here is a way, given a collection of type From, to build with elements of type Elem a collection of type To."

```
final class RNA private (val groups: Array[Int], val length: Int)
  extends IndexedSeq[Base] with IndexedSeqLike[Base, RNA] {

  import RNA._

  // Mandatory re-implementation of `newBuilder` in `IndexedSeq`
  override protected[this] def newBuilder: Builder[Base, RNA] =
    RNA.newBuilder

  // Mandatory implementation of `apply` in `IndexedSeq`
  def apply(idx: Int): Base = {
    if (idx < 0 || length <= idx)
      throw new IndexOutOfBoundsException
    Base.fromInt(groups(idx / N) >> (idx % N * S) & M)
  }

  // Optional re-implementation of foreach,
  // to make it more efficient.
  override def foreach[U](f: Base => U): Unit = {
    var i = 0
```

```

    var b = 0
    while (i < length) {
      b = if (i % N == 0) groups(i / N) else b >>> S
      f(Base.fromInt(b & M))
      i += 1
    }
  }
}

```

**Listing 25.8 - RNA strands class, final version.**

Now the behavior of `map` and `++` on `RNA2` sequences becomes clearer. There is no `CanBuildFrom` instance that creates `RNA2` sequences, so the next best available `CanBuildFrom` was found in the companion object of the inherited trait `IndexedSeq`. That implicit creates `IndexedSeqs`, and that's what you saw when applying `map` to `rna2`.

To address this shortcoming, you need to define an implicit instance of `CanBuildFrom` in the companion object of the `RNA` class. That instance should have `typeCanBuildFrom[RNA, Base, RNA]`. Hence, this instance states that, given an `RNA` strand and a new element type `Base`, you can build another collection which is again an `RNA` strand. Listing 25.8 and Listing 25.9 show the details.

```

object RNA {

  private val S = 2          // number of bits in group
  private val M = (1 << S) - 1 // bitmask to isolate a group
  private val N = 32 / S      // number of groups in an Int

  def fromSeq(buf: Seq[Base]): RNA = {
    val groups = new Array[Int]((buf.length + N - 1) / N)
    for (i <- 0 until buf.length)
      groups(i / N) |= Base.toInt(buf(i)) << (i % N * S)
    new RNA(groups, buf.length)
  }

  def apply(bases: Base*) = fromSeq(bases)

  def newBuilder: Builder[Base, RNA] =
    new ArrayBuffer mapResult fromSeq

  implicit def canBuildFrom: CanBuildFrom[RNA, Base, RNA] =
    new CanBuildFrom[RNA, Base, RNA] {
      def apply(): Builder[Base, RNA] = newBuilder
      def apply(from: RNA): Builder[Base, RNA] = newBuilder
    }
}

```

**Listing 25.9 - RNA companion object—final version.**

Compared to class `RNA2` there are two important differences. First, the `newBuilder` implementation has moved from the `RNA` class to its companion object. The `newBuilder` method in class `RNA` simply forwards to this definition. Second, there is now an implicit `CanBuildFrom` value in object `RNA`. To create such an object you need to define two `apply` methods in the `CanBuildFrom` trait. Both create a new builder for an `RNA` collection, but they differ in their argument list. The `apply()` method simply creates a new builder of the right type. By contrast, the `apply(from)` method takes the original

collection as argument. This can be useful to adapt the dynamic type of builder's return type to be the same as the dynamic type of the receiver. In the case of RNA this does not come into play because RNA is a final class, so any receiver of static type RNA also has RNA as its dynamic type. That's why `apply(from)` also simply calls `newBuilder`, ignoring its argument.

That is it. The RNA class in Listing 25.8 implements all collection methods at their natural types. Its implementation requires a little bit of protocol. In essence, you need to know where to put the `newBuilder` factories and the `canBuildFrom` implicits. On the plus side, with relatively little code you get a large number of methods automatically defined. Also, if you don't intend to do bulk operations like `take`, `drop`, `map`, or `++` on your collection, you can choose to not go the extra length and stop at the implementation shown in Listing 25.6.

The discussion so far centered on the minimal amount of definitions needed to define new sequences with methods that obey certain types. But in practice you might also want to add new functionality to your sequences or override existing methods for better efficiency. An example of this is the overridden `foreach` method in class RNA. `foreach` is an important method in its own right because it implements loops over collections. Furthermore, many other collection methods are implemented in terms of `foreach`. So it makes sense to invest some effort optimizing the method's implementation.

The standard implementation of `foreach` in `IndexedSeq` will simply select every *i*'th element of the collection using `apply`, where *i* ranges from 0 to the collection's length minus one. So this standard implementation selects an array element and unpacks a base from it once for every element in an RNA strand. The overriding `foreach` in class RNA is smarter than that. For every selected array element it immediately applies the given function to all bases contained in it. So the effort for array selection and bit unpacking is much reduced.

### **Integrating new sets and maps**

As a second example you'll learn how to integrate a new kind of map into the collection framework. The idea is to implement a mutable map with `String` as the type of keys by a "Patricia trie".<sup>[2]</sup> The term Patricia is an abbreviation for "Practical Algorithm to Retrieve Information Coded in Alphanumeric." The idea is to store a set or a map as a tree where subsequent characters in a search key determines uniquely a descendant tree.

For instance, a Patricia trie storing the five strings, "abc", "abd", "al", "all", "xy", would look like the tree given in Figure 25.1. To find the node corresponding to the string "abc" in this trie, simply follow the subtree labeled "a", proceed from there to the subtree labeled "b" to finally reach its subtree labeled "c". If the Patricia trie is used as a map, the value that's associated with a key is stored in the nodes that can be reached by the key. If it is a set, you simply store a marker saying that the node is present in the set.

```
import collection._

class PrefixMap[T]
  extends mutable.Map[String, T]
    with mutable.MapLike[String, T, PrefixMap[T]] {
```

```

var suffixes: immutable.Map[Char, PrefixMap[T]] = Map.empty
var value: Option[T] = None

def get(s: String): Option[T] =
  if (s.isEmpty) value
  else suffixes get (s(0)) flatMap (_.get(s substring 1))

def withPrefix(s: String): PrefixMap[T] =
  if (s.isEmpty) this
  else {
    val leading = s(0)
    suffixes get leading match {
      case None =>
        suffixes = suffixes + (leading -> empty)
      case _ =>
    }
    suffixes(leading) withPrefix (s substring 1)
  }

override def update(s: String, elem: T) =
  withPrefix(s).value = Some(elem)

override def remove(s: String): Option[T] =
  if (s.isEmpty) { val prev = value; value = None; prev }
  else suffixes get (s(0)) flatMap (_.remove(s substring 1))

def iterator: Iterator[(String, T)] =
  (for (v <- value.iterator) yield ("", v)) ++
  (for ((chr, m) <- suffixes.iterator;
    (s, v) <- m.iterator) yield (chr ++ s, v))

def += (kv: (String, T)): this.type = { update(kv._1, kv._2); this }

def -= (s: String): this.type = { remove(s); this }

override def empty = new PrefixMap[T]
}

```

**Listing 25.10 - An implementation of prefix maps with Patricia tries.**



**Figure 25.1 - An example Patricia trie.**

Patricia tries support very efficient lookups and updates. Another nice feature is that they support selecting a subcollection by giving a prefix. For instance, in the tree in Figure 25.1 you can obtain the sub-collection of all keys that start with an "a" simply by following the "a" link from the root of the tree.

Based on these ideas we will now walk you through the implementation of a map that's implemented as a Patricia trie. We call the map a `PrefixMap`, which means that it provides a method `withPrefix` that selects a submap of all keys starting with a given prefix.

We'll first define a prefix map with the keys shown in Figure 25.1:

```
scala> val m = PrefixMap("abc" -> 0, "abd" -> 1, "a1" -> 2,
  "a11" -> 3, "xy" -> 4)
m: PrefixMap[Int] = Map((abc,0), (abd,1), (a1,2), (a11,3),
  (xy,4))
```

Then calling `withPrefix` on `m` will yield another prefix map:

```
scala> m withPrefix "a"
res14: PrefixMap[Int] = Map((bc,0), (bd,1), (1,2), (11,3))
```

Listing 25.10 shows the definition of `PrefixMap`. This class is parameterized with the type of associated values `T`, and extends `mutable.Map[String, T]` and `mutable.MapLike[String, T, PrefixMap[T]]`. You have seen this pattern already for sequences in the RNA strand example. Then as now inheriting an implementation class such as `MapLike` serves to get the right result type for transformations such as `filter`.

A prefix map node has two mutable fields: `suffixes` and `value`. The `value` field contains an optional value that's associated with the node. It is initialized to `None`. The `suffixes` field contains a map from characters to `PrefixMap` values. It is initialized to the empty map. You might ask, Why did we pick an immutable map as the implementation type for `suffixes`? Would not a mutable map been more standard since `PrefixMap` as a whole is also mutable? The answer is that immutable maps that contain only a few elements are very efficient in both space and execution time.

For instance, maps that contain fewer than 5 elements are represented as a single object. By contrast, as described in Section 17.2, the standard mutable map is a `HashMap`, which typically occupies around 80 bytes, even if it is empty. So if small collections are common, it's better to pick immutable over mutable. In the case of Patricia tries, we'd expect that most nodes, except the ones at the very top of the tree, would contain only a few successors. So storing these successors in an immutable map is likely to be more efficient.

Now have a look at the first method that needs to be implemented for a map: `get`. The algorithm is as follows: To get the value associated with the empty string in a prefix map, simply select the optional value stored in the root of the tree. Otherwise, if the key string is not empty, try to select the submap corresponding to the first character of the string. If that yields a map, follow up by looking up the remainder of the key string after its first character in that map. If the selection fails, the key is not stored in the map, so return with `None`. The combined selection over an option value is elegantly expressed using `flatMap`. When applied to an optional value, `ov`, and a closure, `f`, which in turn returns an optional value, `ov flatMap f` will succeed if both `ov` and `f` return a defined value. Otherwise `ov flatMap f` will return `None`.

The next two methods to implement for a mutable map are `+=` and `-=`. In the implementation of Listing 25.10, these are defined in terms of two other methods: `update` and `remove`. The `remove` method is very similar to `get`, except that before returning any associated value, the field containing that value is set to `None`. The `update` method first calls `withPrefix` to navigate to the tree node that needs to be updated, then sets the `value` field of that node to the given value. The `withPrefix` method navigates through the tree, creating sub-maps as necessary if some prefix of characters is not yet contained as a path in the tree.

The last abstract method to implement for a mutable map is `iterator`. This method needs to produce an iterator that yields all key/value pairs stored in the map. For any given prefix map this iterator is composed of the following parts: First, if the map contains a defined value, `Some(x)`, in the `value` field at its root, then `("", x)` is the first element returned from the iterator. Furthermore, the iterator needs to traverse the iterators of all submaps stored in the `suffixes` field, but it needs to add a character in front of every key string returned by those iterators. More precisely, if `m` is the submap reached from the root through a character `chr`, and `(s, v)` is an element returned from `m.iterator`, then the root's iterator will return `(chr +: s, v)` instead.

This logic is implemented quite concisely as a concatenation of two for expressions in the implementation of the `iterator` method in Listing 25.10. The first for expression iterates over `value.iterator`. This makes use of the fact that `Option` values define an iterator method that returns either no element, if the option value is `None`, or exactly one element `x`, if the option value is `Some(x)`.

```
import scala.collection.mutable.{Builder, MapBuilder}
import scala.collection.generic.CanBuildFrom

object PrefixMap {
  def empty[T] = new PrefixMap[T]

  def apply[T](kvs: (String, T)*): PrefixMap[T] = {
    val m: PrefixMap[T] = empty
    for (kv <- kvs) m += kv
    m
  }

  def newBuilder[T]: Builder[(String, T), PrefixMap[T]] =
    new MapBuilder[String, T, PrefixMap[T]](empty)

  implicit def canBuildFrom[T]
    : CanBuildFrom[PrefixMap[_], (String, T), PrefixMap[T]] =
    new CanBuildFrom[PrefixMap[_], (String, T), PrefixMap[T]] {
      def apply(from: PrefixMap[_]) = newBuilder[T]
      def apply() = newBuilder[T]
    }
}
```

**Listing 25.11 - The companion object for prefix maps.**

Note that there is no `newBuilder` method defined in `PrefixMap`. There is no need because maps and sets come with default builders, which are instances of class `MapBuilder`. For a mutable map the default builder starts with an empty map and then adds successive elements using the map's `+=` method.



Mutable sets behave the same. The default builders for immutable maps and sets use the non-destructive element addition method `+`, instead of method `+=`. However, in all these cases, to build the right kind of set or map, you need to start with an empty set or map of this kind. This is provided by the `empty` method, which is the last method defined in `PrefixMap`. In Listing 25.10, this method simply returns a fresh `PrefixMap`.

We'll now turn to the companion object `PrefixMap`, which is shown in Listing 25.11. In fact it is not strictly necessary to define this companion object, as class `PrefixMap` can stand well on its own. The main purpose of object `PrefixMap` is to define some convenience factory methods. It also defines a `CanBuildFrom` implicit to make typing work out better.

The two convenience methods are `empty` and `apply`. The same methods are present for all other collections in Scala's collection framework so it makes sense to define them here too. With the two methods, you can write `PrefixMap` literals like you do for any other collection:

```
scala> PrefixMap("hello" -> 5, "hi" -> 2)
res0: PrefixMap[Int] = Map((hello,5), (hi,2))

scala> PrefixMap.empty[String]
res2: PrefixMap[String] = Map()
```

The other member in object `PrefixMap` is an implicit `CanBuildFrom` instance. It has the same purpose as the `CanBuildFrom` definition in the last section: to make methods like `map` return the best possible type. For instance, consider mapping a function over the key/value pairs of a `PrefixMap`. As long as that function produces pairs of strings and some second type, the result collection will again be a `PrefixMap`. Here's an example:

```
scala> res0 map { case (k, v) => (k + "!", "x" * v) }
res8: PrefixMap[String] = Map((hello!,xxxxx), (hi!,xx))
```

The given function argument takes the key/value bindings of the prefix map `res0` and produces pairs of strings. The result of the `map` is a `PrefixMap`, this time with value type `String` instead of `Int`. Without the `canBuildFrom` implicit in `PrefixMap` the result would just have been a general mutable map, not a prefix map.

## Summary

If you want to fully integrate a new collection class into the framework, you need to pay attention to the following points:

1. Decide whether the collection should be mutable or immutable.
2. Pick the right base traits for the collection.
3. Inherit from the right implementation trait to implement most collection operations.
4. If you want `map` and similar operations to return instances of your collection type, provide an implicit `CanBuildFrom` in your class's companion object.

## 25.4 CONCLUSION

You have now seen how Scala's collections are built and how you can build new kinds of collections. Because of Scala's rich support for abstraction, each new collection type can have a large number of methods without having to reimplement them all over again.

### Footnotes for Chapter 25:

[1] Ideally, everything should be defined in one place only, but there are a few exceptions where things needed to be redefined.

[2] Morrison, "PATRICIA—Practical Algorithm To Retrieve Information Coded in Alphanumeric" [Mor68]